

---

# Plant Traits: Leveraging pre-trained models

---

**Sean Huh**

School of Computer Science  
University of Waterloo  
Waterloo, ON, N2L 3G1  
s5huh@uwaterloo.ca  
report due: August 12

## Abstract

Please note: This project makes a critical failing by using leaked test data in the pre-trained SWin model. Despite this, the effort will be to demonstrate the model's effectiveness as if the SWin model was pre-trained on solely training data.

The basis of this project's model was exploring how to learn from multi-modal data with limited compute resources. Tuning an image transformer is resource intensive and so use of a pre-tuned SWin transformer was used for image feature extraction. Other methods beyond concatenation for combination of image and tabular features were explored but ultimately deemed to overfit to the training data. The final model's hyperparameters were fine-tuned with random searching. This resulted in an  $R^2 \approx 0.55$ . Re-testing using a safe pre-tuned SWin is the next step to test this model's true efficacy. See [github](#) for code.

## 1 Introduction

The problem is to predict 6 plant trait labels given photographs from the iNaturalist database and plant trait data. The training set has 43363 images with 164 columns of tabular data. Architecture of the model used to make predictions is given in Figure 1.

In brief, the methodology used was to leverage a pre-trained SWin model that had been tuned on the PlantsTraits2024 competition images as an image feature extractor. A simpler network was used on the tabular data to extract tabular features. These features were concatenated and processed through multiple layers to predict the 6 plant traits.

Predominantly, the question was how to balance effectiveness with efficiency. Training vision models are computationally heavy and impractical to train from ground-up. Thus the use of pre-trained vision models was necessary for this problem, utilizing models that could already capture the features of the plant images effectively.

There was a decision over how to merge the tabular data and image features (extracted from the pre-trained model). Preliminary exploration on predicting with just image features resulted in decent performance. To fully train a more accurate model required combining the multi-modal data in some way.

## 2 Related Works

The skeleton of the code was adapted from Kaggle user's HdJoJo's methodology for tuning a model on PlantsTraits2024 using just image data. Fine-tuning pre-trained SWin models for various datasets (ADE20K, etc.) is quite effective as seen in Liu et al. (2021). There is an inherent feature representation in the pre-trained SWin for general images (i.e. ImageNet) that is then modified slightly for

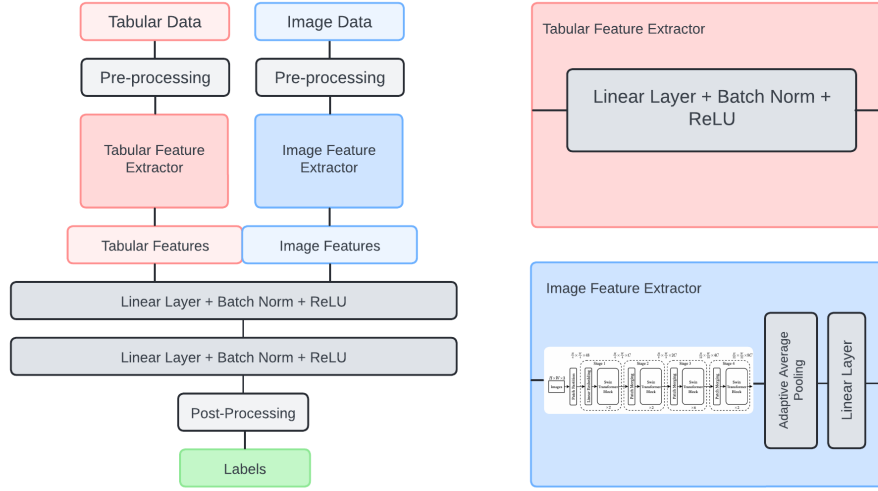


Figure 1: This multi-modal model is used to predict plant traits given tabular and image data.

specific sets of images. This idea is extended here for this model where the desire is to extract the inherent feature representation for plant images.

Ensemble models train individual models to predict the labels and then weigh these models to create a larger model. In contrast, multi-modal models combine features across modalities by 'fusing' them and then training on the fusion to predict the labels. Kim et al. (2021) showcases the efficiency of one such fusion method, ViLT. This method incorporates a language embedding into the vision transformer to learn context from the text. Inspired by this and to a limited capacity (due to having obtained the aforementioned vision transformer), trials involving cross-attention between the tabular and image features were attempted.

Regarding SWin, Liu2021's description of architecture indicates each channel can be used to represent some high-level 2D representation of a feature of the original image. Adaptive average pooling (with stride 1) condenses the output of SWin into a one-dimensional feature vector.

### 3 Main Results

Following Shi et al. (2021), the test labels were augmented by log-transforming and normalizing. Due to the high variance of the test label values, the training data was clipped within a 0.005 and 0.985 percentile. The tabular data was also augmented by log-transforming any traits with high skewness ( $>1$ ) then normalizing. This was done to handle heavily-skewed distributions and reduce scale sensitivity in preparation of training. Figure 2 demonstrates the effect both these augmentations have on the distribution of the six traits of interest and six randomly chosen features from the tabular data.

The images were augmented by applying the same augmentations used by HdJoJo during pre-training, then upsampled to  $(382 \times 382)$  to match the SWin model's input dimensions. Applying random flips and contrasts along with image cropping was done to improve the model's generalization and reduce overfitting.

After augmenting the data, the model feeds the image and tabular data into feature extractors respectively. The tabular feature extractor is a simple neural network that is trained along with the rest of the model during training. The image feature extractor is a pre-trained SWin model tuned on the PlantsTraits2024 training images. The initial model predicted the final 6 traits directly, so the extractor removes the last prediction layer and pools each channel into a single float, creating a feature vector. This is then fed into a down-resolution neural network to condense these features further. Most importantly regarding use of the pre-trained SWin model, all layers are frozen in place during training except this last down-resolution network. Because the initial model was trained on a

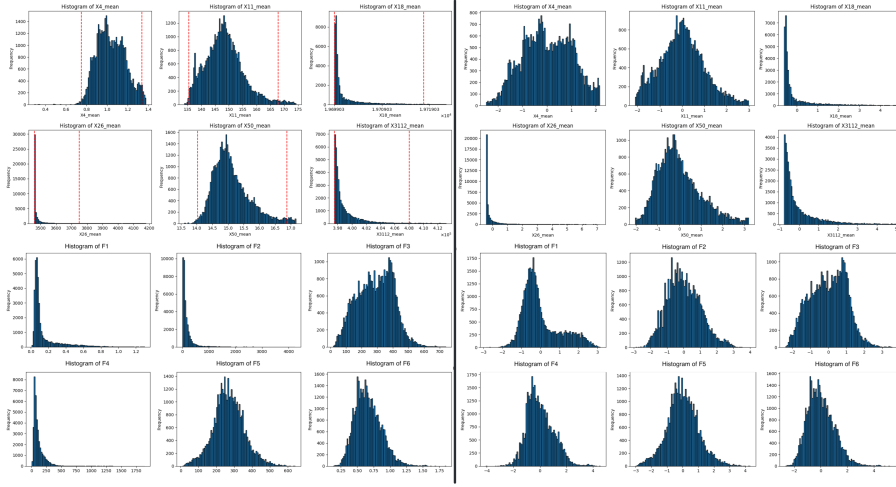


Figure 2: By augmenting the labels and features, performance improves due to normalization and reduced scale sensitivity.

66 similar dataset and to save computation time, the assumption is that SWin’s output is already tuned  
 67 well and does not require further training. In ablation testing, this is proven to be correct when the  
 68 SWin model is used directly on the dataset with decent accuracy, see Table 1. Figure 1 shows the  
 69 components of the image feature extractor and tabular feature extractor.

70 After both the image and tabular features are extracted, they are then fused into a single layer by  
 71 concatenating them. This is then brought through two more layers with batch normalization and  
 72 ReLU as usual in-between.

73 As a trial, cross-attention mechanisms were also explored as fusion methods. Ideally, the tabular  
 74 and image features would attend to one another to capture the most relevant information (per head)  
 75 using the other modality as context. These contextualized features would then be concatenated and  
 76 used to predict the final traits using a final layer. The model (as depicted in Figure 3) was deployed  
 77 and trained over three epochs. Ultimately, the model showed a degradation in performance versus  
 simple concatenation as a fusion method.

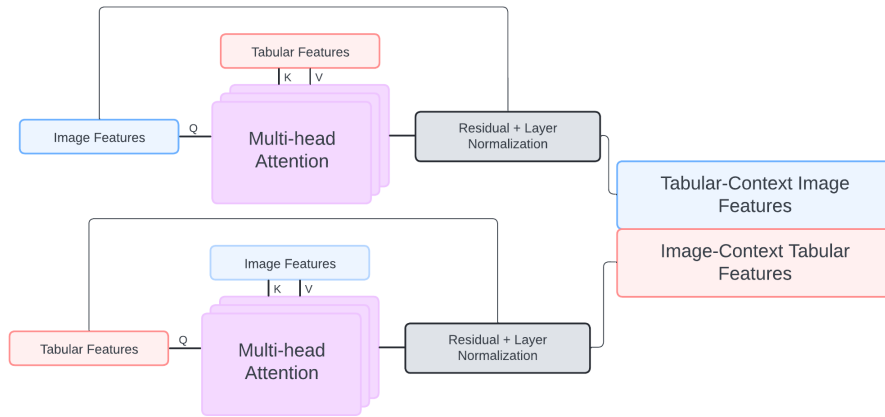


Figure 3: Alternative methods of fusion proved unsuccessful.

Table 1: Model Comparison on stand-alone SWin and Hyperparameter Tuning.  
In order: [Learn Rate, Weight Decay, Image Features, Tab. Features, Combo. L1, Combo. L2].

Model	Loss average	MAE average	R <sup>2</sup>
SWin-standalone	-	-	0.29
1e-3/1e-3/128/256/512/1024/128	0.207	0.469	0.47
1e-3/1e-3/256/64/256/2048/128	0.206	0.468	0.474
1e-2/1e-4/128/128/512/512/256	0.205	0.469	0.473
1e-2/1e-4/512/512/512/512/256	<b>0.203</b>	<b>0.466</b>	<b>0.481</b>

The model was trained using a OneCycleLR scheduler with AdamW. By increasing the learning rate at the beginning of the training, the model is able to converge quickly to optimal values, reducing need for extended training time as describe in Smith (2017). The model is trained using SmoothL1Loss due to concern over outliers disrupting the model’s convergence. Pseudocode for how the gradient and learning rate is updated is available in Algorithm 1.

**Algorithm 1:** AdamW with OneCycleLR using SmoothL1Loss

```

1 for step = 0, 1, 2, ... (num_epochs · n) do
2    $\hat{y} \leftarrow \text{model}(X)$ 
3    $\frac{\partial l}{\partial \hat{y}} \leftarrow \begin{cases} y_i - \hat{y}_i & \text{if } |y_i - \hat{y}_i| \leq -1, \\ \text{sign}(y_i - \hat{y}_i) & \text{otherwise} \end{cases} \quad // \text{ } i=1, 2, \dots, 6$ 
4    $\hat{g} \leftarrow \frac{\partial l}{\partial W} = \frac{\partial l}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial W} \quad // \text{ for some arbitrary layer weights } W$ 
5    $lr \leftarrow$ 
84    $\begin{cases} lr_{min} + \frac{1}{2}(lr_{max} - lr_{min})(1 + \cos\left(\frac{\text{step} - \text{pct} \cdot \text{num\_epochs} \cdot n}{(1 - \text{pct}) \cdot \text{num\_epochs} \cdot n} \cdot \pi\right)) & \text{if } \frac{\text{step}}{\text{num\_epochs} \cdot n} > \text{pct}, \\ lr_{max} + \frac{1}{2}(lr_{min} - lr_{max})(1 + \cos\left(\frac{\text{step}}{\text{pct} \cdot \text{num\_epochs} \cdot n} \cdot \pi\right)) & \text{otherwise} \end{cases}$ 
6    $m_{\text{step}} \leftarrow \beta_1 \cdot m_{\text{step}-1} + (1 - \beta_1) \cdot \hat{g}$ 
7    $v_{\text{step}} \leftarrow \beta_2 \cdot v_{\text{step}-1} + (1 - \beta_2) \cdot \hat{g}^2$ 
8    $m_{\text{step}} \leftarrow m_{\text{step}} / (1 - \beta_1)$ 
9    $v_{\text{step}} \leftarrow v_{\text{step}} / (1 - \beta_2)$ 
10   $W \leftarrow W - lr \cdot m_{\text{step}} / (\sqrt{v_{\text{step}}} + \epsilon) - lr \cdot \text{weight\_decay} \cdot W$ 

```

Several hyperparameters were tuned by performing three epochs of training using different configurations via Random Search, then testing against the test set. This was chosen over a Box Method due to limited compute resources. Results for the standalone Swin-model along with the 5 most-diverse hyperparameter configurations tested are summarized in Table 1.

The final training occurred over 6 epochs of training and dimensions using hyperparameters [1e-4, 1e-2, 512, 512, 512, 512, 256]. Running this model three times resulted in an  $R^2$  standard deviation of 0.002256.

## 4 Conclusion

This model uses the idea of multi-part training for multi-modal models. By using a SWin model pre-tuned to the image dataset, the surrounding model was relatively simple and efficient. Limited computation resources may do not allow multiple training sessions of complex models, so it is beneficial to train a complex model per modality once, and then be able to repurpose for other tasks. Despite the pre-tuned model using leaked test data, the model should still work with at least some effectiveness had the model been tuned on strictly training data.

## 99 Acknowledgement

100 Thank you to TA Saber Malekmohammadi for helping me in the creation of this report. I made a  
101 mistake in using a model that was pre-trained on leaked test data. After I explained the mistake, he  
102 was supportive in understanding what had occurred and allowed me to still submit my model.

103 Thank you to professor Yao-Liang Yu for being foundational to my understanding of many of the  
104 topics and techniques used for this project.

105 Figure 1 adopts from Liu et al. (2021).

## 106 References

107 Kim, W., B. Son, and I. Kim (2021). “ViLT: Vision-and-Language Transformer Without Convolution  
108 or Region Supervision”. *arXiv: 2102.03334 [stat.ML]*.

109 Liu, Z., Y. Lin, Y. Cao, H. Hu, Y. Wei, Z. Zhang, S. Lin, and B. Guo (2021). “Swin Transformer:  
110 Hierarchical Vision Transformer using Shifted Windows”. *arXiv: 2103.14030 [cs.CV]*.

111 Shi, L., Q. Wei, L. Xie, W. Wang, and Y. Zheng (2021). “Growing deep learning-based image seg-  
112 mentation for bone age assessment: A systematic review and future perspectives”. *Scientific Re-*  
113 *ports*, vol. 11, no. 1, p. 14553.

114 Smith, L. N. (2017). “Super-Convergence: Very Fast Training of Neural Networks Using Large  
115 Learning Rates”. *arXiv preprint arXiv:1708.07120*.